

Predicting Complaint Eligibility for Accountability Mechanisms

Accountability Counsel is a nonprofit that supports communities whose human rights, livelihoods, and environments are threatened by development projects financed by powerful investors and banks. Part of this work involves helping communities file complaints with accountability mechanisms, which determine whether a complaint meets eligibility requirements before review and investigation. In partnership with the UChicago Data Science Clinic, this project explored whether a large language model system could identify eligibility issues in draft complaints and generate feedback to help communities strengthen submissions.

The team analyzed historical complaint filings to examine eligibility patterns and test whether an LLM could assist with eligibility prediction. Because the dataset was small and imbalanced, results were mixed, especially when the model tried to identify the relatively limited number of ineligible cases. The filings also show that submission volume and eligibility outcomes vary over time.

For the CAO mechanism, the model frequently labeled complaints as eligible, even when cases were actually ineligible. This suggests the current approach is better at flagging potentially eligible complaints than screening out likely ineligible complaints. Despite these limitations, the project established a reproducible pipeline for organizing complaint data, evaluating model performance, and identifying areas for improvement, including dataset expansion, prompt refinement, and fine-tuning or other model adaptation.




Figure 1. Complaint eligibility outcomes for CAO filings by year. The figure shows how the number of eligible and ineligible submissions changes over time, indicating that complaint volume and outcomes are not constant across periods.




Figure 2. Classification matrix for CAO eligibility predictions. The model frequently predicts “Eligible,” including for several complaints that were actually ineligible, which suggests that the current prompting approach is overly lenient.