

Data, Policy & Innovation Centre

Introduction

The Government of Odisha received hundreds of thousands of citizen grievances each year, many accompanied by supporting documents that provided additional details about complaints. However, the unstructured format, inconsistent scan quality, and multilingual nature of these materials made large-scale analysis difficult. In collaboration with the Data, Policy & Innovation Centre, the team developed a system to convert these documents into structured, privacy-preserving text that could be analyzed more efficiently.

Pipeline

The team built a document classification model that predicted page attributes such as language, and whether the text was handwritten or typed. Using these predictions, the team filtered for English documents, extracted their text, and applied a redaction process that removed personally identifiable information such as names, phone numbers, and addresses.

During the winter quarter, the team processed a larger dataset containing more than 76,000 documents and over 300,000 pages. All pages were classified and stored in a centralized document repository, with approximately 30 percent identified as English. Text was extracted from roughly 96,000 English pages and sensitive information was removed. The team also developed models that summarized grievances and grouped complaints into categories like infrastructure, healthcare, and public services,, helping reveal patterns across submissions.




Figure above: complete project statistics