The Government of Odisha in India receives hundreds of thousands of citizen grievances each year, many accompanied by supporting documents with additional complaint details. The unstructured format, inconsistent scan quality, and multilingual nature of these materials had made large-scale analysis difficult. In collaboration with the Data, Policy & Innovation Centre (DPIC), the team developed an automated grievance redressal system to convert these documents into structured, privacy-preserving text for more efficient review by government officials.

The team built a processing pipeline beginning with a model that identified each document's language and whether its text was handwritten or typed. Since the team could only feasibly verify English results, roughly 30% of documents were selected for further processing. Text was then extracted and scrubbed of personally identifiable information, including names, phone numbers, and addresses, to protect citizen privacy. A third model identified each page by document type, filtering out pages unlikely to aid summarization. Two final models worked in tandem: one produced concise grievance summaries, and another assigned category labels to route complaints to the appropriate government department. The completed pipeline was delivered to DPIC with usage demonstrations and retraining guidance, providing a functional proof-of-concept ready for further deployment.

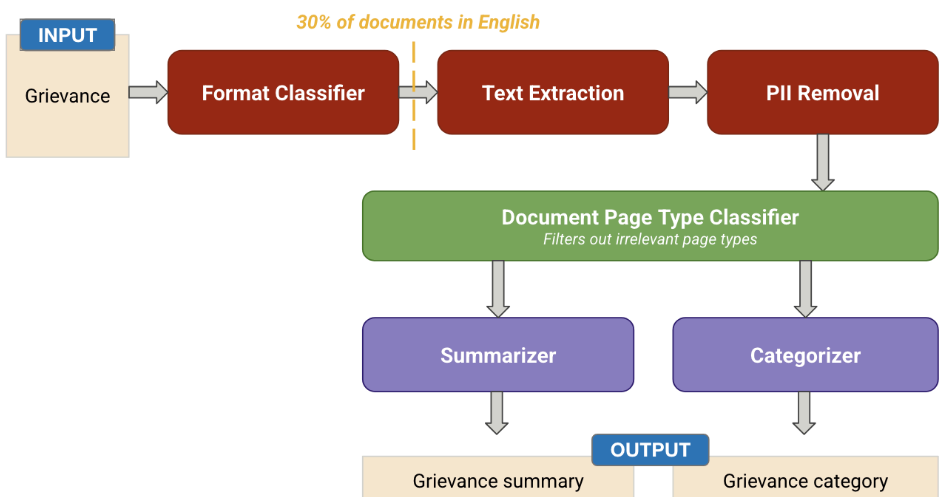

Figure 1. *The five-stage processing pipeline. Raw citizen grievance documents are classified by format and stripped of personally identifiable information before a document-type classifier removes pages unlikely to aid summarization. The document is finally summarized and categorized, and routed to the appropriate government department.*